

Labeling Questions Asked On GitHub Issue Trackers

Aidin Rasti

Electrical Engineering and Computer Science

University of Ottawa

Ottawa, Canada

arast040@uottawa.ca

Abstract—One of the issues faced by the maintainers of popular open source software is the triage of newly reported issues. Many of the issues submitted to issue trackers are questions. Many people ask questions on issue trackers about their problem instead of using a proper QA website like StackOverflow. This may seem insignificant but for many of the big projects with thousands of users, this leads to spamming of the issue tracker. Reading and labeling these unrelated issues manually is a serious time consuming task and these unrelated questions add to the burden. In fact, most often maintainers demand to not submit questions in the issue tracker.

To address this problem, first, we leveraged dozens of patterns to clean text of issues, we removed noises like logs, stack traces, environment variables, error messages, etc. Second, we have implemented a classification-based approach to automatically label unrelated questions. Empirical evaluations on a dataset of more than 102,000 records show that our approach can label questions with an accuracy of over 81%.

Index Terms—issue tracking system, natural language processing, machine learning, mining software repositories

I. INTRODUCTION

With the prevalence of open source software, authors of projects kindly carry on the responsibility of supporting users and providing documentation. The relation between users and authors is a two-way relation, developers rely on feedback and reports from users to improve their software. These feedback reports come in the form of bug reports, questions, features, suggestions, and enhancements, or what we technically call "issue"s. Most developers prefer to use the issue tracking software just as a means for developing the software, it is preferred that questions related to the workings or documentation of the software be directed to forums or QA websites. And by questions we mean everything that is, asking about how to fix an error/problem, asking about features and documentation, asking for help, etc.

These unrelated questions only add to the clutter of issues trackers, especially in bigger projects. Developers have to put a tremendous amount of effort to triage or close issues. Providing a simple automated tool to detect potential unrelated questions can help project managers and software developers to focus on more practical issues.

Other works of software defect classification focus on classifying bugs, features, enhancements, etc. Most of them don't even include questions. Since most of them are based on the data provided from internal issue trackers of projects

Testing with counter_cache and fixtures #19000

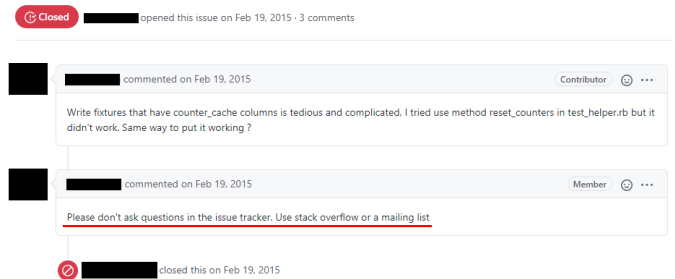


Fig. 1. A question submitted to the issue tracker of rails project on GitHub.

like JBoss, Apache Foundation, they are trained on a structured and clean dataset. However, for developers that use the built-in issue tracker of platforms like GitHub, the problem still remains. Moreover, we want to focus on implementing a binary classifier to filter questions, instead of the multi-class classifiers that other works propose. Maybe we can say that the goals are a little different or even complementary. For example, we can first filter spam by using our proposed classifier then when we are more confident that the reported issue is related, we can use another classifier to automatically categorize it. This may improve the accuracy of automatic categorization too. Also, this enables deployment of those classifiers on a broader range of projects.

In this report, we want to investigate the feasibility of labeling unrelated questions in the issue tracker of popular open source software on platforms like GitHub. We want to evaluate the performance of different classification algorithms. And finally, to compare the performance of a binary classification (question, not-question) approach to other multi-class methods used in other works.

As mentioned earlier, our goal is to train a binary classifier to automatically label questions. We have used a previously available dataset of GitHub issues provided by [1]. Cleaning and pre-processing the dataset narrowed down the count of our available records to 102,198. The pre-processing part was the hardest part of our implementation since the text of issues usually contain machine-generated texts like logs and stack traces. After extracting the text part of issues, we have used state of the art sentence embedding techniques to convert text

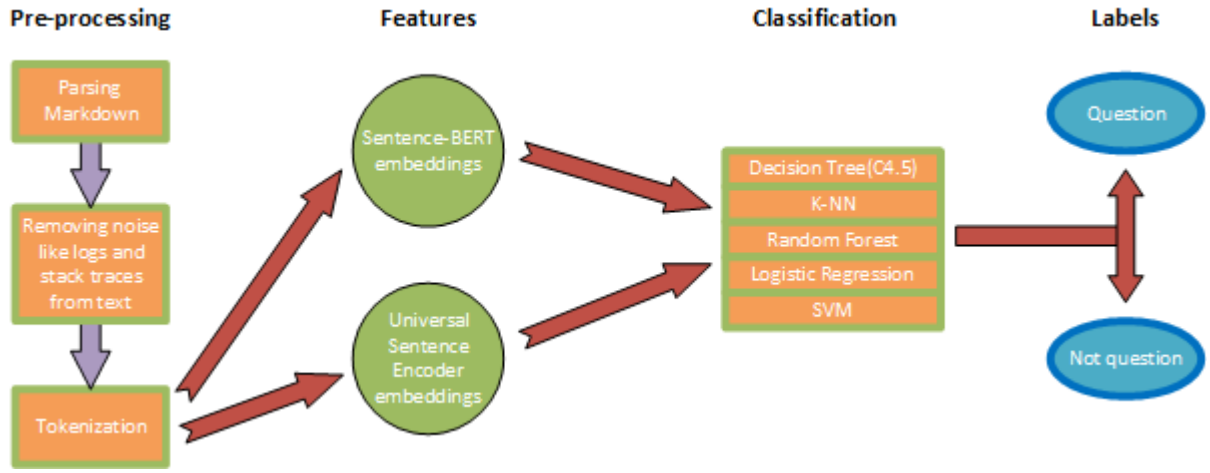


Fig. 2. Summary of our approach at a high level.

data to numerical vectors. By evaluating several classification algorithms, we achieved the best result with the Logistic Regression algorithm.

First, In Section II we discuss other related works. In Section III we describe our dataset, data pre-processing methods, and technical implementation. In Section IV we show the performance of our implementation and discuss the results of different classification methods. Finally, Section VI concludes this report.

II. RELATED WORKS

The work of Antoniol et al. [2] is one of the defect classification works. They trained classifiers with Naive Bayes, ADTree, and Logistic Regression on the text data of 1,800 issues extracted from Mozilla, Eclipse, and JBoss. Their work is focused on classifying whether a submitted issue is bug or non-bug (enhancement, feature request, etc). Zhou et al. [3] improve on [2] by including other structured attributes available in an issue tracker, such as assignee, dates, severity, into the machine learning model. Pingclasai et al. [4] build a classifier using LDA topic modeling for reclassification of issues. They used Decision Tree, Naive Bayes classifier, and Logistic Regression to classify issues to bug and not-bug. Another work by Cubranic and Murphy [5] use Naive Bayes for text classification to find an assignee for a bug report. Hanmin and Xin used an LSTM neural network [6] to classify issues to bug and non-bugs too, they showed that with an LSTM neural network they can achieve better performance than some other machine learning based approaches.

Kochhar et al. [7] address the problem of misclassified reports in issue trackers. They have extracted features from 7,000 issues from five open source projects to classify them into 13 categories. Unlike the other mentioned works, Kochhar et al. build a multi-class classifier using textual and structural data from issues. But their categories do not include question.

To the best of our knowledge, we couldn't find any issue classification works which included questions. Furthermore,

almost all of the mentioned works were done on a small dataset (less than 8,000 records) collected from internal issue trackers.

III. APPROACH

In summary, our approach involves three main phases:

- 1) The first phase is cleaning and pre-processing the text data of issues. We used around 400 regular expressions to remove noise from the text of issues. Most of the issues are not properly formatted and they contain stack traces, log lines, code snippets, command lines, environment variables, configurations, IP addresses, identifiers, flags, timestamps, etc. We have also applied the standard NLP pre-processing tasks such as tokenization and removing punctuations.
- 2) The second phase is computing a sentence embedding of each text document provided by the previous step. We have used the Sentence Bert [8] technique to convert our documents to a high dimensional vector representation. The Sentence Bert is one the best performing sentence embedding techniques according to benchmarks [8].
- 3) The final phase is training a classifier on the extracted features. We have evaluated the performance of SVM, Decision Tree (C4.5), Logistic Regression, Random Forest, and k-NN classification algorithms.

In the following subsections, we explain our dataset, pre-processing methods, a few background definitions, and our implementation.

A. Dataset

We used the RapidRelease [1] dataset. It contains 2 million issues from active and popular GitHub repositories of 18 programming languages. The dataset is a SQL database. The labels of each issue, which were applied by the project maintainers, are also available in this dataset.

With a simple query, we extracted frequently applied labels. You can see the count of the top 10 labels in Table I. We have selected these top labels, bug, duplicate, enhancement,

TABLE I
TOP 10 APPLIED LABELS

Label	Count
bug	65735
enhancement	43673
question	34248
cla: yes	17476
duplicate	12210
feature	9235
stale	7996
documentation	7814
invalid	7658
cla signed	7306

wontfix, feature, improvement, and question to find more labels based on them. For example, some projects may use the label "type: bug" to represent a bug. Also, many of the issues have more than one label. The labels are stored in the database as string in a comma-separated format. For example, an issue with both labels of bug and duplicate is stored as "bug,duplicate". For each of these base labels, we queried issues that contained their strings. For example, for features, we looked for labels that contained the string "feature". We also queried their usage count and extracted labels for each category based on their usage count. For example, The top five applied labels that contained the "feature" string are: "feature", "feature request", "kind/feature", "cat:feature", "type: feature". A label with a higher usage indicates that it is a relevant label for the intended purpose. This query resulted in thousands of records for each of the base labels, however, many of them were only applied a handful of times. Therefore we have only selected labels from these lists that at least were used 50 times. For our classifier, we have two categories of not-question and question. We used the base labels bug, duplicate, enhancement, wontfix, feature, improvement, and their similar labels for the not-question category. We used the question or similar labels for the question category. Given that many issues have more than one label, we removed issues that had labels like "bug,question" or "support,bug" from the list of frequent labels. We used this list of labels to extract issues. The list of labels used can be found in our code repository in the labels folder for each one of the base labels. With this process we extracted 46,549 issues for the question category and 280,829 for the not-question.

B. Date cleaning and pre-processing

After extracting issues the pre-processing phase begins. GitHub issues are written in the Markdown format. We used the marked [9] library to parse and compile issues, this library compiles Markdown to HTML. Following compiling issues to HTML we extracted the text part of issues from HTML. Also, during the translation to HTML, we removed image, table, code, pre, and header tags. If the author of an issue had used the Markdown properly then after removing these tags only the text part of an issue would have remained. However, many developers do not submit issues in the proper format and we still have to further clean texts.

We applied more than 400 regular expressions to remove the remaining noise from the text of issues. Specifically, we have removed dependency trees, emojis, code snippets, comments, logs, error message, stack traces, timestamp, date times, command lines, environment variables, identifiers, HTML, markup tags, module versions, IP addresses, emails, GitHub user handles, URIs, file paths, and digits. These regular expressions were designed by manually skimming through texts. Suppose we want to extract stack traces, most of the stack traces contain several consecutive lines that begin with the "at" or "in" word or they contain the file extension of that specific programming language (".js", ".py"). Also, stack traces of a platform like Java or .NET are very similar to each other and creating a pattern to extract them is easy. By manually analyzing our regular expression for stack traces we identified thousands of occurrences. As another example, suppose we want to remove log lines. Many developers copy logs or debug output of the software to help diagnose the problem. Almost all of the logs contain a timestamp on each line, therefore, if we see several consecutive lines that contain timestamps, there is a high probability that they are software generated log outputs.

By manually applying our regular expressions and inspecting dozens of results we have fine tuned them to not catch incorrect patterns. Of course, these regular expressions are not perfect and there are still a handful of incorrect matches in a large dataset, but given that they can match hundreds of correct patterns, we opted to apply them to remove noise from the text of issues. We have also removed several punctuations except those that indicated the end of a sentence (such as .,;!?). The whole list of all used regular expressions and the subscript is available in the code repository.

The next step of pre-processing was filtering issues that were not in the English language. We used the langdetect [10] library to determine the language of each issue and then filter the ones that were not English.

After cleaning the dataset we concatenated the title and body of issues then we used the nltk library for tokenization. We used its default tokenizer which is TreebankWordTokenizer. Following tokenization, we marked issues that had more than five and lower than 200 tokens for training. This reduced the records of the question category to 42,198 and the not-question category to 262,425. Also, for the question category we only selected issues that were closed.

C. Sentence Embeddings

To understand the concept of sentence embeddings we should first understand word embeddings. Word embeddings are essentially vector representations of words. There are a few techniques to train numerical vector representation of words. Word2vec [11], [12] and GloVe [13] are the most famous word embeddings algorithms. The advantage of word embeddings over older techniques like One-hot encoding and Bag-of-words is that semantically similar words appear close to each other in the vector space and the relation among the words are preserved. One popular example of this case is the similarity

between the words “king” and “queen”, for example, we can calculate a vector close to the vector of word “queen” by calculating the formula:

$$\text{king} - \text{men} + \text{woman} = \text{queen} \quad (1)$$

Word embeddings give us a numerical vector for each word. Sentence embeddings are an extension to the word embeddings technique. Sentence embeddings are vector representations of a sentence, paragraph, or even a whole text document. One naive way of calculating the sentence embeddings of a text or sentence is to average its word embeddings, however, this method does not perform well for NLP tasks. Sentence-BERT [8] is a state of the art algorithm to derive sentence embeddings. Sentence-BERT builds on top of the BERT [14] networks to calculate meaningful sentence embeddings. Another algorithm to calculate sentence embeddings is Universal Sentence Encoder [15]. Vectors generated by both of the two mentioned sentence embedding algorithms can be compared using the cosine similarity function. Sentence embeddings perform extremely well for similarity and classification tasks. Since our task is classification, we decided to use these sentence embeddings to compute a numerical vector for the text of each issue in our cleaned dataset. We computed embeddings using both Universal Sentence Encoder and Sentence-BERT then we trained classifiers for both of them to compare results. The input to these models is a string and the output is a numerical vector. We used the `roberta-large-nli-stsb-mean-tokens` pre-trained model for Sentence-BERT which is trained over the SNLI [16] and Multi NLI [17] datasets. This model generates 1024 dimension vectors. We used the `universal-sentence-encoder-large-v5` pre-trained model of Universal Sentence Encoder algorithm. This model generates 512 dimension vectors. This model is trained over a variety of unsupervised data sources from the web. They used data from Wikipedia and question-answer websites. The resulting labeled vectors of both of these embeddings are available in our code repository.

D. Training classifiers

After extracting feature vectors we now have a dataset ready to train classification models. We have two datasets, one generated from Universal Sentence Encoder and one generated from Sentence-BERT. The Universal Sentence Encoder records have 512 dimensions while the Sentence-BERT ones have 1024 dimensions. We used the Weka UI toolkit [18] to train classifiers. We have trained classifiers using Logistic Regression, Decision Tree (C4.5), Random Forest, and Support Vector Machine. In this section, we briefly review each of the classification algorithms.

- **Logistic Regression:** Logistic Regression can be used for binary classification. Instead of predicting a continuous value like Linear Regression, it predicts the probability belonging to two different classes. During the training process, it finds a logistic function to calculate this probability.

- **Decision Tree (C4.5):** The C4.5 decision tree algorithm calculates the information gain ratio for all of the attributes then splits data at each node by the attribute with the highest normalized information gain ratio. It recursively continues to split data until the tree is complete.
- **Random Forest:** Random Forest is an Ensemble Learning algorithm. Basically, it consists of several uncorrelated decision trees, each used for classification. The class with the majority of votes becomes the final results.
- **Support Vector Machine:** SVM is one of the best performing linear classification algorithms. It can classify data by finding the best hyperplane that separates the dataset into two given classes. To allow room for generalization it also selects the hyperplane which has the most margin from both categories. It can also work on a non-linearly separable dataset by using a kernel function.
- **k-Nearest Neighbors:** To use the k-NN algorithm for classification the raw dataset becomes the model, there is no training process. When a record is queried for its class, the algorithm finds the k closest data points in the raw dataset according to a distance function (e.g. Euclidean distance). The class of the queried record is determined by the majority vote of its neighbors. The number of neighbor points (k) to use for classification is an input parameter, the best k value can be determined by evaluating results over the training split.

IV. EMPIRICAL EVALUATION

We have implemented our approach to answer our three research questions. Our questions are:

- RQ1) *Is it possible to prevent spamming in issue trackers of popular open source software? is it possible to label questions asked in issue trackers?*
- RQ2) *Which classification algorithms yield goods result for this specific task?*
- RQ3) *Is the binary classification (question, not-question) a good approach? does it help or improve results?*

A. RQ1 and RQ2

In order to answer our first two questions, we have evaluated the performance of several popular classification algorithms over both datasets. One dataset containing feature vectors generated by the Sentence-BERT algorithm and another generated by the Universal Sentence Encoder algorithm. We have two labels for classification. The `question` category has 42,198 issues. For the `not-question` category, we randomly chose 60,000 issues from a total of 262,425. This made our training dataset more balanced while not breaking the 60-40 threshold. In total, our dataset contains 102,198 records. The exact same issues were used for both of the sentence embedding algorithms, however, the generated feature vectors have different dimensions, vectors of the Sentence-BERT algorithm has 1024 dimension while vectors of the Universal Sentence Encoder has 512 dimensions. Vectors generated by both these algorithms are normalized and we did not perform any normalization on them for any of the

TABLE II
ACCURACY OF CLASSIFICATION ALGORITHMS TRAINED OVER BOTH
EMBEDDINGS WITH A 70-30 SPLIT.

Sentence Embedding	Classifier	Accuracy
Universal Sentence Encoder (512 features)	k-NN	78.38%
	Decision Tree(C4.5)	68.12%
	Logistic Regression	81.68%
	Random Forest	79.05%
	SVM	78.18%
Sentence-BERT (1024 features)	k-NN	69.88%
	Decision Tree(C4.5)	60.41%
	Logistic Regression	78.42%
	Random Forest	71.35%
	SVM	79.86%

classification algorithms. We used the Weka toolkit to train classification models. If Weka toolkit normalized values by default for any of the classification algorithms, we turned it off before training models.

We trained classifiers with the Logistic Regression, Decision Tree (C4.5), Random Forest, and Support Vector Machine algorithms. For all of the classifiers, we ran the training process with their default parameters in Weka. You can find these parameters in the code repository under the `weka` folder. We trained all of the algorithms over 70% of the dataset and we evaluated them over the 30% test split. Also, we trained each of these algorithms two times since we had two datasets. Due to the limited computation resources and time, we were not able to perform cross-validation. Still, a 30% test split over a dataset with more than 100,000 records is large and diverse enough to give reliable results. The output of Weka for each algorithm is available in our code repository too.

In Table II you can find the accuracy of each classifier. Logistic Regression has the highest accuracy among Universal Sentence Encoder data and SVM has the highest accuracy among Sentence-BERT. The best classifier is the Logistic Regression trained over Universal Sentence Encoder embeddings with an accuracy of 81.68%. It almost performs two percent better than the best classifier trained over the Sentence-BERT embeddings. In addition, we can observe that classifiers trained over the Universal Sentence Encoder embeddings have more consistent results. Almost all of its classifiers perform better with an accuracy of at least 78% (except the decision tree).

To evaluate our approach better we examine other metrics too, accuracy alone is not a good metric for evaluation. We have gathered the confusion matrix of each classifier in Fig. 3. By looking at it we see that the hardest part is labeling actual questions, all of the algorithms almost struggle at classifying them. Logistic Regression over Universal Sentence Encoder embeddings, which had the highest accuracy, has the highest TP rate for the question category too (74.8%). SVM over Universal Sentence Encoder embeddings has the highest true negatives category. But The SVM classifier also has the lowest false positive numbers, the Logistic Regression has almost 1000 more false positives. In classifiers trained over the Sentence-BERT embeddings, SVM has both the highest TP (69.8%) and the lowest FP rate (13.2%). But the Random

Forest classifier has the highest true negatives with a little margin higher than the SVM.

To further compare the results, we have provided Precision, Recall, and F-measure rates in Table III. The Logistic algorithm trained over Universal Sentence Encoder embeddings has the highest F-measure rates. In the Sentence-BERT embeddings category, the SVM classifier has the highest F-measure. By analyzing the precision rates we can still see that Logistic Regression trained with universal Sentence Encoder embeddings still has the best performance, its average precision is higher than any other classifier. Although, as we can see in Fig. 3 it has much higher false positives than the SVM and Random Forest. It may be more sensible to use these two instead of the Logistic Classifier.

Regarding RQ1, given the evaluations and empirical evidence, we can conclude that it is feasible to classify questions in issue trackers of platforms like GitHub with pretty good precision. We saw that we can achieve an accuracy rate of 81.6%. To answer RQ2 we trained several models with popular algorithms. We assessed two different Sentence Embedding techniques to find which is more suitable for our task. We observed that the Universal Sentence Encoder has better and more consistent performance for our task. The answer to RQ2 is not a single algorithm, if we only measure the accuracy metric, the Logistic Regression works best. However, if we analyze other metrics we may choose other algorithms. For example, the Random Forest model trained with Universal Sentence Encoder embeddings has an accuracy of 79.00% but it has much lower false positives than the Logistic Regression classifier.

We noticed that the bottleneck in the accuracy of almost all of the classifiers is categorizing issues that are actually questions. They can label issues that are not questions with

TABLE III
PRECISION AND RECALL RATES OF

Classifier	Class	Precision	Recall	F-Measure
<i>Universal Sentence Encoder</i>				
k-NN	NQ	79.2%	85.8%	82.3%
	Q	77.0%	67.9%	72.1%
Decision Tree(C4.5)	NQ	73.0%	72.6%	72.8%
	Q	61.3%	61.8%	61.5%
Logistic Regression	NQ	83.0%	86.5%	84.7%
	Q	79.6%	74.8%	77.1%
Random Forest	NQ	77.3%	91.2%	83.6%
	Q	83.1%	61.8%	70.9%
SVM	NQ	76.2%	91.4%	83.1%
	Q	82.9%	59.4%	69.2%
<i>Sentence-BERT</i>				
k-NN	NQ	78.0%	68.3%	72.8%
	Q	61.2%	72.2%	66.2%
Decision Tree(C4.5)	NQ	66.7%	66.0%	66.3%
	Q	51.6%	52.3%	51.9%
Logistic Regression	NQ	80.1%	84.4%	82.2%
	Q	75.6%	69.8%	72.6%
Random Forest	NQ	71.1%	86.9%	78.2%
	Q	72.1%	48.9%	58.3%
SVM	NQ	80.6%	86.8%	83.6%
	Q	78.6%	69.8%	73.9%

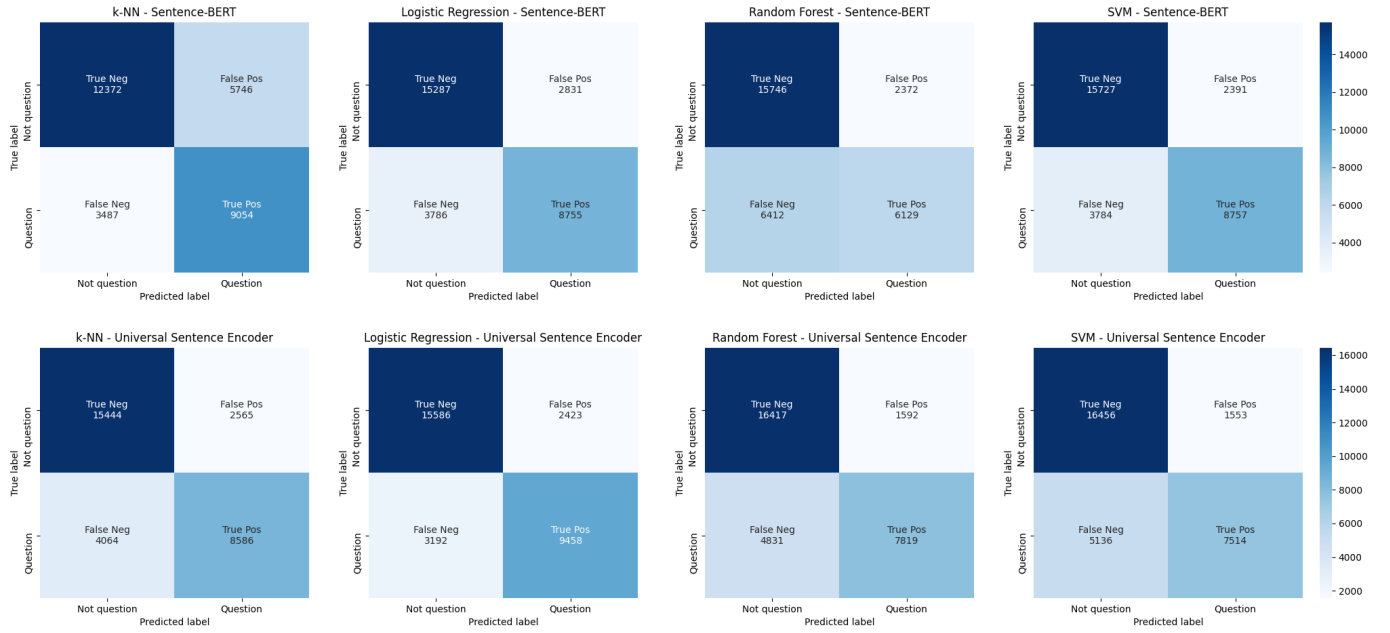


Fig. 3. Heat map of confusion matrix of trained models over both sentence embedding techniques.

good accuracy. This problem can be attributed to the fact that many of the submitted questions in issues trackers will be converted to another type of issue by the maintainers. For example, a submitted question about a feature may interest developers and they decide to implement it, thus they will apply a “feature” label to the issue which is asking a question. The same happens with other labels such as “bug” too, a submitted question about an error may in fact turn out to be a bug. Since our dataset is not manually labeled for the purpose of our task and we depend on the labels provided by authors of projects this problem can’t be addressed trivially. Manually labeling issues may help address this problem but it is time consuming. Another approach may be to analyze issues with an unsupervised algorithm.

B. RQ3

With RQ3 we are interested in comparing our approach with other works. However, we found out a direct comparison with other works is not possible. As we saw in Section II, almost all of the works in the defect classification area are multi-class classification problems. Nevertheless, we think that our approach and other works are complementary. We can apply both of them in a pipeline fashion. First, we can separately label questions submitted to an issue tracker, then apply another categorization method to predict the class of defect.

C. Limitations and Further Improvements

Due to the size of our dataset and limited computation power, we were not able to tune hyperparameters of classification algorithms used in Weka. We used the default parameters set by Weka. However, when we searched for the best value

for a few of the parameters over a small sample dataset, we saw a modest improvement in results.

Our initial solution to solve this problem was to train an LSTM network with Word Embeddings. Using lower level embeddings with a neural network may improve results. But we don’t know how much improvement it can yield. As explained in Section III after pre-processing the RapidRelease dataset, we only used issues with more than five and less than 200 tokens to train models. Also, vectors generated by Word Embedding algorithms can have at least 100 dimensions. With this high amount of input and training parameters, training a neural network failed on our device.

As mention above, we filtered issues with long texts, we have only trained models over issues that had lower than 200 tokens. We removed longer issues because the performance of sentence embedding algorithms degrades with longer texts. Still, 200 tokens are long enough for most of the issues and we also counted tokens after removing all the noise (log lines, stack traces) from the text of issue.

D. Threats to Validity

Like all evaluations, our tests are also prone to external validity and generalization. We have only used data from the GitHub issue tracker. But we think that the number of questions asked on the issue tracker of projects that don’t use GitHub (for example, ASF software have their own issue trackers) is very low. Furthermore, the RapidRelease dataset is diverse and includes issues from several developer communities and programming languages.

Another factor is that we relied on the labels applied by developers of projects. The labels may not completely reflect our intended two categories, especially for questions. As we discussed in Subsection IV-C this may be the reason

why classifiers had less precision for the question category. Overall, we tried to use labels that were frequently applied by developers to issues.

V. LESSONS LEARNED

The most important lesson I learned was to not rely on a single solution. It is always best to evaluate several approaches to find the best possible fit. Another thing I noticed is that the accuracy metric alone is not enough to gauge performance. We have to look at other metrics too and analyze more details. The accuracy metric can be misleading.

VI. CONCLUSION

Our goal was analyzing feasibility of detecting unrelated questions submitted to issue trackers in platforms like GitHub. To address this problem We used a previously provided dataset of GitHub issues by the RapidRelease work. We filtered and pre-processed the raw dataset and narrowed count of issues for analysis to approximately 102,000 issues. We were able to filter noise, such as stack traces, logs, etc, from text of issues. To prepare our textual data for machine learning algorithms, we used two of state-of-the-art sentence embeddings techniques, Sentence-BERT and Universal Sentence Encoder, to extract numerical feature vectors. Finally, we trained classifiers over our dataset with five famous classification algorithms. The best result was yielded with the Logistic Regression algorithm trained over Universal Sentence Encoder embeddings with the accuracy rate of 81.68%. Overall, we conclude that it is possible to classify questions submitted to a public issues tracker with a good probability.

REFERENCES

- [1] S. D. Joshi and S. Chimalakonda, "Rapidrelease - a dataset of projects and issues on github with rapid releases," in *2019 IEEE/ACM 16th International Conference on Mining Software Repositories (MSR)*, 2019, pp. 587–591.
- [2] G. Antoniol, K. Ayari, M. Di Penta, F. Khomh, and Y.-G. Guéhéneuc, "Is it a bug or an enhancement? a text-based approach to classify change requests," in *Proceedings of the 2008 Conference of the Center for Advanced Studies on Collaborative Research: Meeting of Minds*, ser. CASCOS '08. New York, NY, USA: Association for Computing Machinery, 2008. [Online]. Available: <https://doi.org/10.1145/1463788.1463819>
- [3] Y. Zhou, Y. Tong, R. Gu, and H. Gall, "Combining text mining and data mining for bug report classification," in *2014 IEEE International Conference on Software Maintenance and Evolution*, 2014, pp. 311–320.
- [4] N. Pingclasai, H. Hata, and K. Matsumoto, "Classifying bug reports to bugs and other requests using topic modeling," in *2013 20th Asia-Pacific Software Engineering Conference (APSEC)*, vol. 2, 2013, pp. 13–18.
- [5] D. Cubranic and G. C. Murphy, "Automatic bug triage using text categorization," in *Software Engineering and Knowledge Engineering*, 2004, p. 92–97.
- [6] H. Qin and X. Sun, "Classifying bug reports into bugs and non-bugs using lstm," in *Proceedings of the Tenth Asia-Pacific Symposium on Internetware*, ser. Internetware '18. New York, NY, USA: Association for Computing Machinery, 2018. [Online]. Available: <https://doi-org.proxy.bib.uottawa.ca/10.1145/3275219.3275239>
- [7] P. S. Kochhar, F. Thung, and D. Lo, "Automatic fine-grained issue report reclassification," in *2014 19th International Conference on Engineering of Complex Computer Systems*, 2014, pp. 126–135.
- [8] N. Reimers and I. Gurevych, "Sentence-bert: Sentence embeddings using siamese bert-networks," in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 11 2019. [Online]. Available: <https://arxiv.org/abs/1908.10084>
- [9] C. Jeffrey, "Marked," 2020. [Online]. Available: <https://marked.js.org/>
- [10] M. Danilak, "langdetect," 2020. [Online]. Available: <https://github.com/Mimino666/langdetect/>
- [11] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient estimation of word representations in vector space," 2013.
- [12] T. Mikolov, I. Sutskever, K. Chen, G. Corrado, and J. Dean, "Distributed representations of words and phrases and their compositionality," 2013.
- [13] J. Pennington, R. Socher, and C. D. Manning, "Glove: Global vectors for word representation," in *Empirical Methods in Natural Language Processing (EMNLP)*, 2014, pp. 1532–1543. [Online]. Available: <http://www.aclweb.org/anthology/D14-1162>
- [14] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "Bert: Pre-training of deep bidirectional transformers for language understanding," 2019.
- [15] D. Cer, Y. Yang, S. yi Kong, N. Hua, N. Limtiaco, R. S. John, N. Constant, M. Guajardo-Cespedes, S. Yuan, C. Tar, Y.-H. Sung, B. Strope, and R. Kurzweil, "Universal sentence encoder," 2018.
- [16] S. R. Bowman, G. Angeli, C. Potts, and C. D. Manning, "A large annotated corpus for learning natural language inference," in *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, 2015.
- [17] A. Williams, N. Nangia, and S. Bowman, "A broad-coverage challenge corpus for sentence understanding through inference," in *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*. Association for Computational Linguistics, 2018, pp. 1112–1122. [Online]. Available: <http://aclweb.org/anthology/N18-1101>
- [18] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. H. Witten, "The weka data mining software: An update," *SIGKDD Explor. Newsl.*, vol. 11, no. 1, p. 10–18, Nov. 2009. [Online]. Available: <https://doi.org/10.1145/1656274.1656278>